

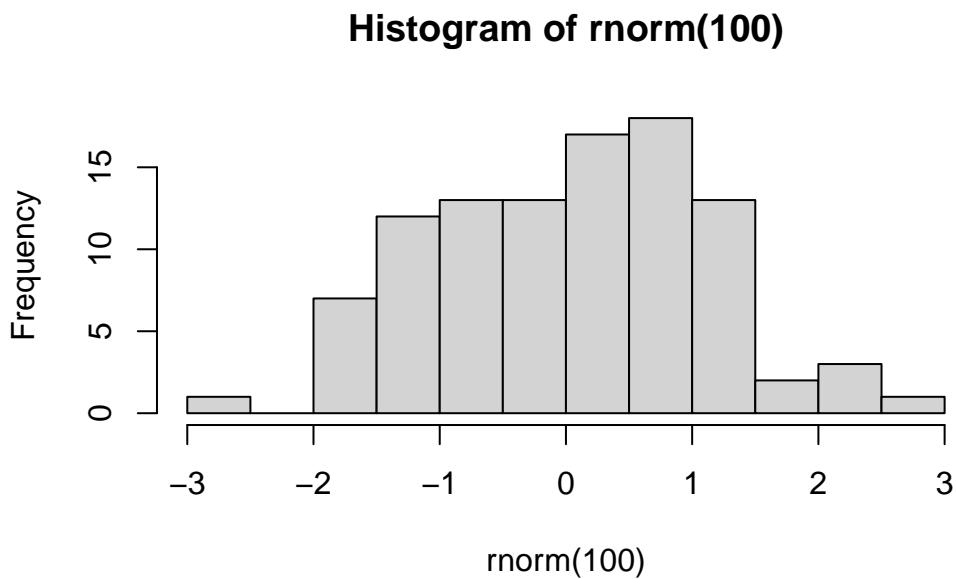
Class07: Machine learning

Bobbie Morales A15443382

Today we will begin our exploration of some “classical” machine learning approaches. We will start with clustering:

Let's first make up some data to cluster where we know what the answer should be.

```
hist(rnorm(100))
```



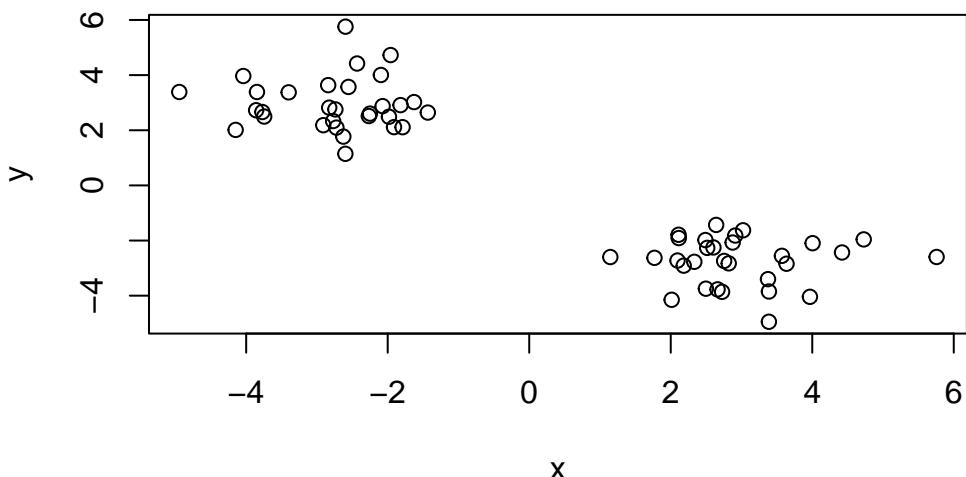
```
x <- c(rnorm(30, mean=-3), rnorm(30, mean=3))
y <- rev(x)

x<- cbind(x, y)
head(x)
```

```
x           y
[1,] -4.148901 2.013092
[2,] -2.725030 2.095612
[3,] -2.071963 2.875000
[4,] -3.848006 3.387797
[5,] -2.598403 1.146123
[6,] -2.266136 2.515443
```

a wee peak at x with `plot()`

```
plot(x)
```



The main function in “base” R for K-means clustering is called `kmeans()`

```
k <- kmeans(x, centers = 2)
```

Q. How big are the clusters (i.e their size)

```
k$size
```

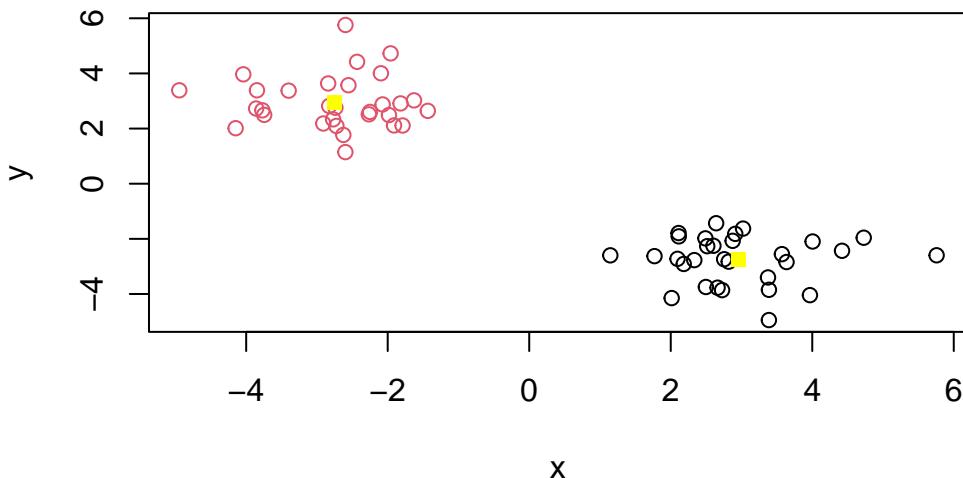
```
[1] 30 30
```

Q. What clusters do my data points reside in?

k\$cluster

Q. Make a plot of our data colored by cluster assignment - ie. make a result figure

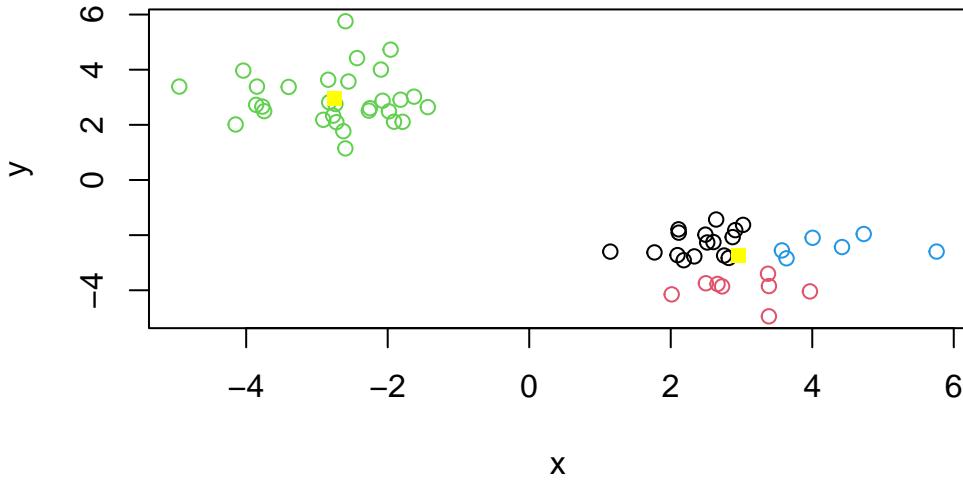
```
plot(x,col=k$cluster)
points(k$centers, col = "yellow", pch=15)
```



Q Cluster with k-means into 4 clusters and plot your results above.

```
k4 <- kmeans(x, centers = 4)

plot(x,col=k4$cluster)
points(k$centers, col = "yellow", pch=15)
```



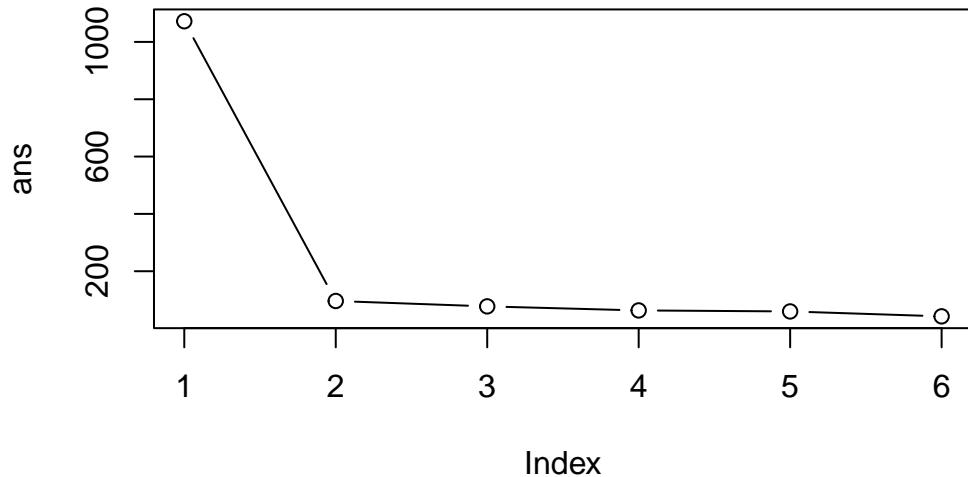
Q. Run kmeans with values center(i.e values of k) equal to 1 to 6.

```
k1 <- kmeans(x, centers=1)$tot.withinss
k2 <- kmeans(x, centers=2)$tot.withinss
k3 <- kmeans(x, centers=3)$tot.withinss
k4 <- kmeans(x, centers=4)$tot.withinss
k5 <- kmeans(x, centers=5)$tot.withinss
k6 <- kmeans(x, centers=6)$tot.withinss

ans <- c(k1, k2, k3, k4, k5, k6)
```

Or use a for loop

```
ans <- NULL
for(i in 1:6) {
  ans <- c(ans, kmeans(x, centers=i)$tot.withinss)
}
plot(ans, type = "b")
```



Hierarchical Clustering

The main function in “base” R for this is called `hclust()`

```
d <- dist(x)
hc <- hclust(d)
hc
```

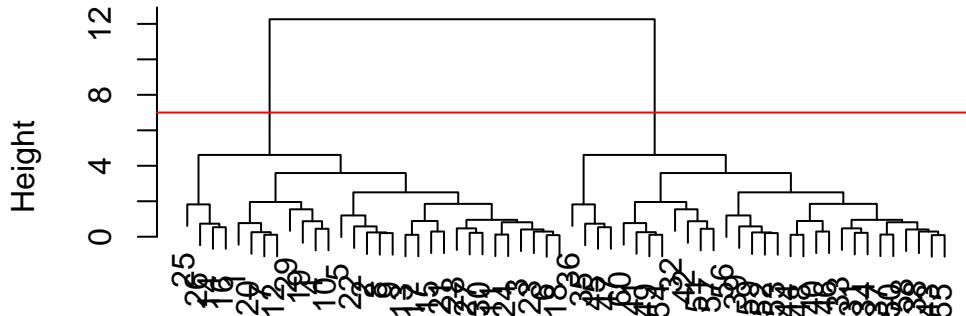
```
Call:
hclust(d = d)

Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

```
#hclust(x)
```

```
plot(hc)
abline(h=7, col="red")
```

Cluster Dendrogram

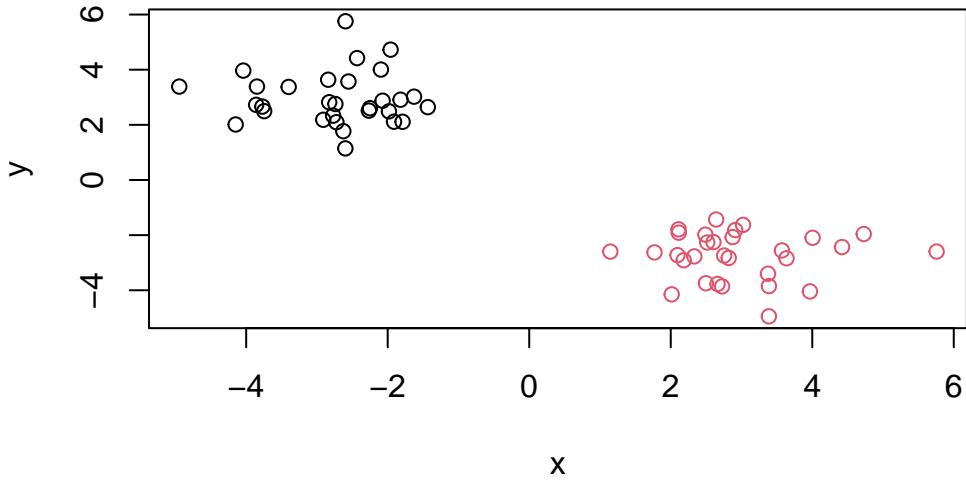


```
d  
hclust (*, "complete")
```

obtain clusters from `hclust` result object `hc` we “cut” the tree to yield different sub branches.
For this we use the `cutree()` function

```
grps <- cutree(hc, h=7)
```

```
plot(x, col=grps)
```



Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

```
dim(x)
```

```
[1] 17 5
```

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
dim(x)
```

```
[1] 17 4
```

```
x <- read.csv(url, row.names=1)
head(x)
```

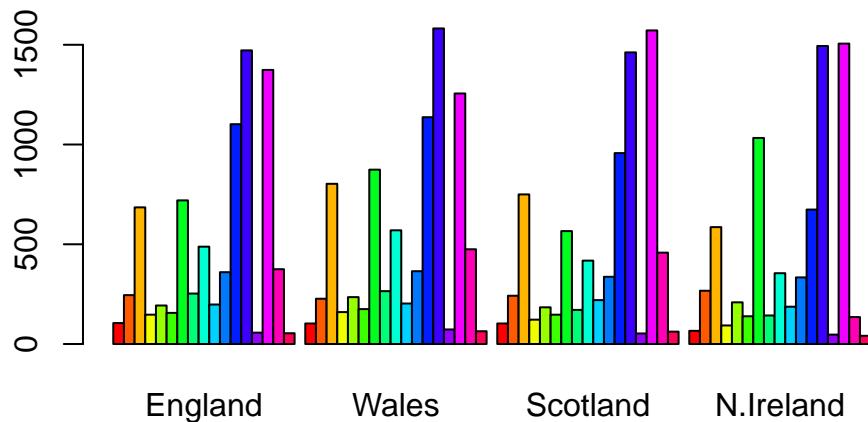
	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

The second approach is more robust. It reads the row names correctly

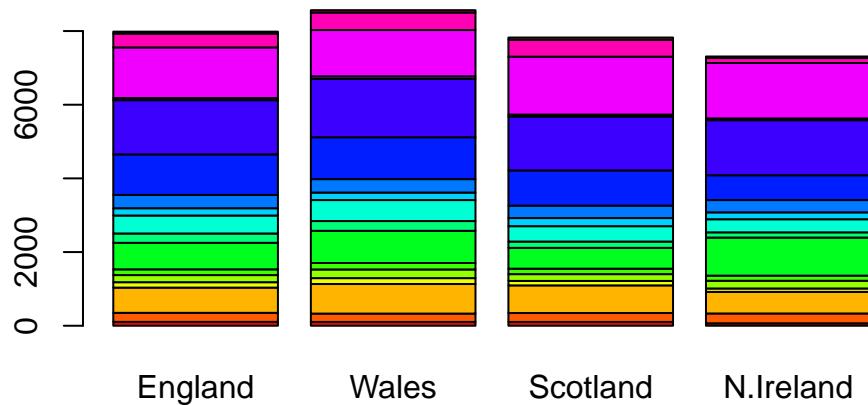
using base R

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above barplot() function results in the following plot?

```
barplot(as.matrix(x), col=rainbow(nrow(x)))
```



Removing the beside resulted in the plot.

Q4: Changing what optional argument in the above ggplot() code results in a stacked barplot figure?

```
library(tidyr)

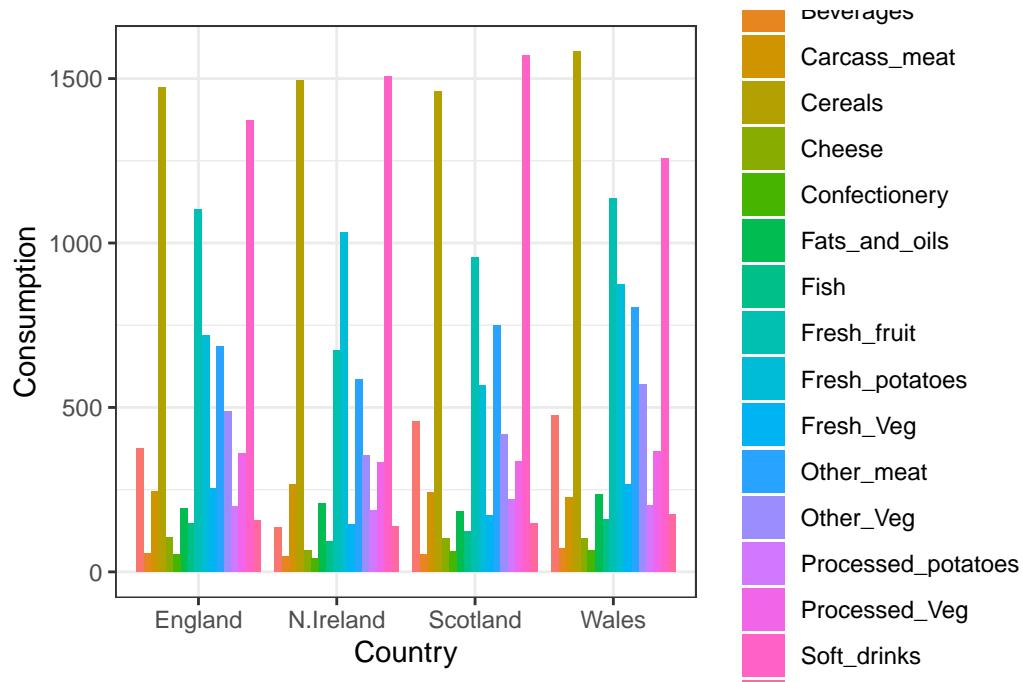
# Convert data to long format for ggplot with `pivot_longer()`
x_long <- x |>
  tibble::rownames_to_column("Food") |>
  pivot_longer(cols = -Food,
                names_to = "Country",
                values_to = "Consumption")

dim(x_long)
```

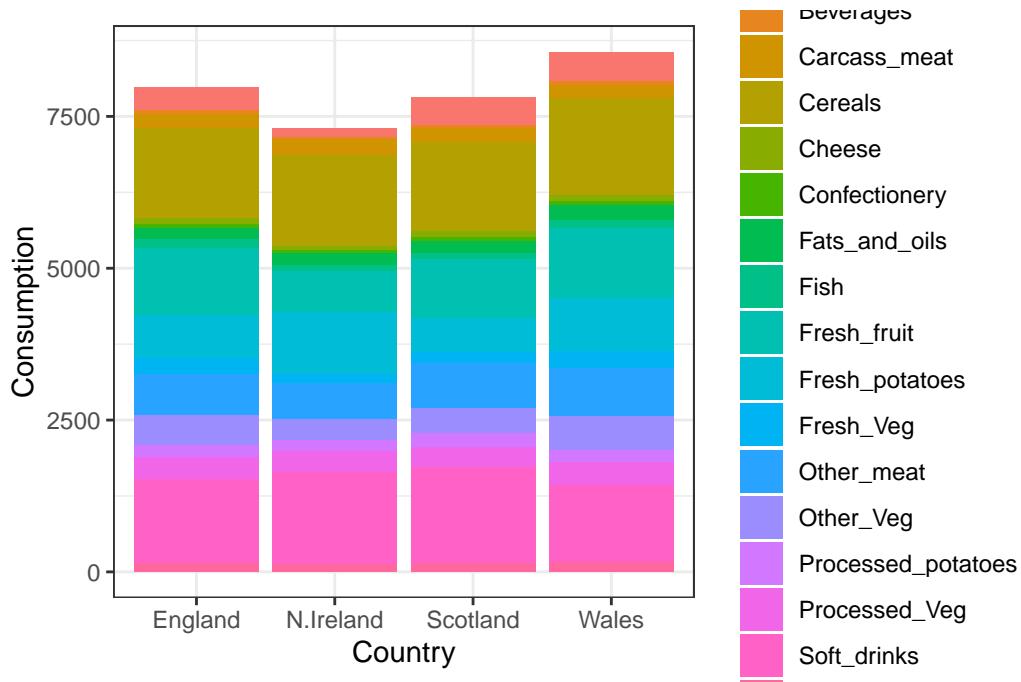
```
[1] 68  3
```

```
library(ggplot2)

ggplot(x_long) +
  aes(x = Country, y = Consumption, fill = Food) +
  geom_col(position = "dodge") +
  theme_bw()
```



```
ggplot(x_long) +
  aes(x = Country, y = Consumption, fill = Food) +
  geom_col() +
  theme_bw()
```

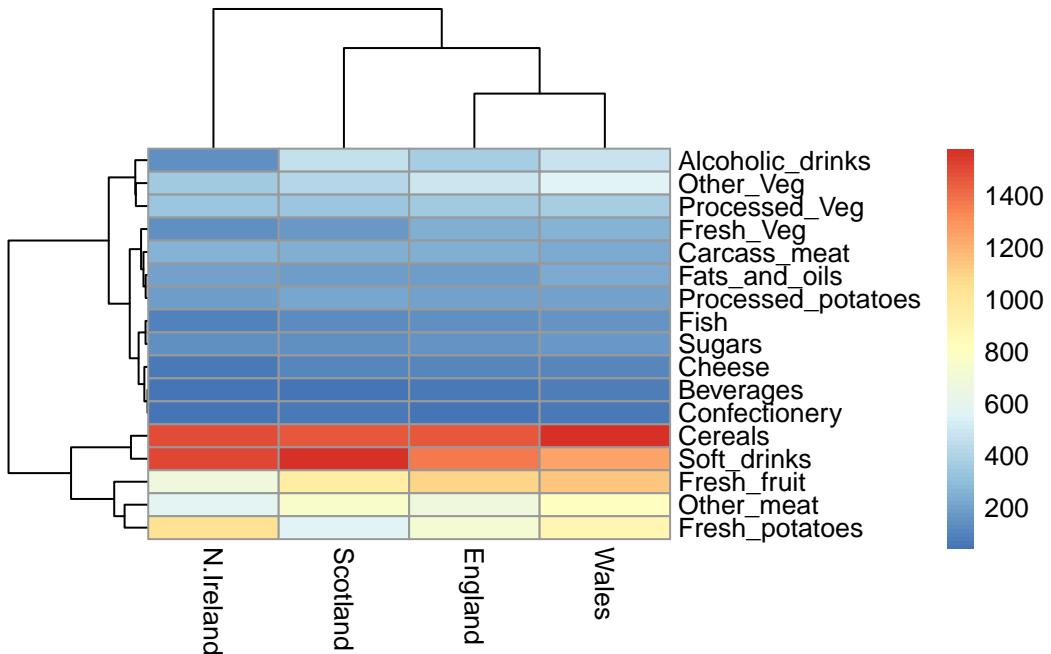


Removing the position argument made the plot.

Q5: We can use the pairs() function to generate all pairwise plots for our countries. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

It is difficult to interpret the plot because the axes are not clear to me. It is taking the four countries we are examining and putting them in a display. They plot each country against another country for comparison. Each dot corresponds to the foods we are looking at. Foods along the diagonal mean there are similar quantities of that food consumed in those 2 countries.

```
library(pheatmap)
pheatmap( as.matrix(x) )
```



Q6. Based on the pairs and heatmap figures, which countries cluster together and what does this suggest about their food consumption patterns? Can you easily tell what the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Wales and England look the most similar in their food consumption. It is difficult to read this plot.

PCA to the rescue

The main function in “base” R for PCA is called `prcomp()`.

As we want to do PCA on the food data for the different countries we will want the foods in the columns (taking the transpose of the data)

```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Our result object is called `pca` and it has a `.$x` component that we will look at first

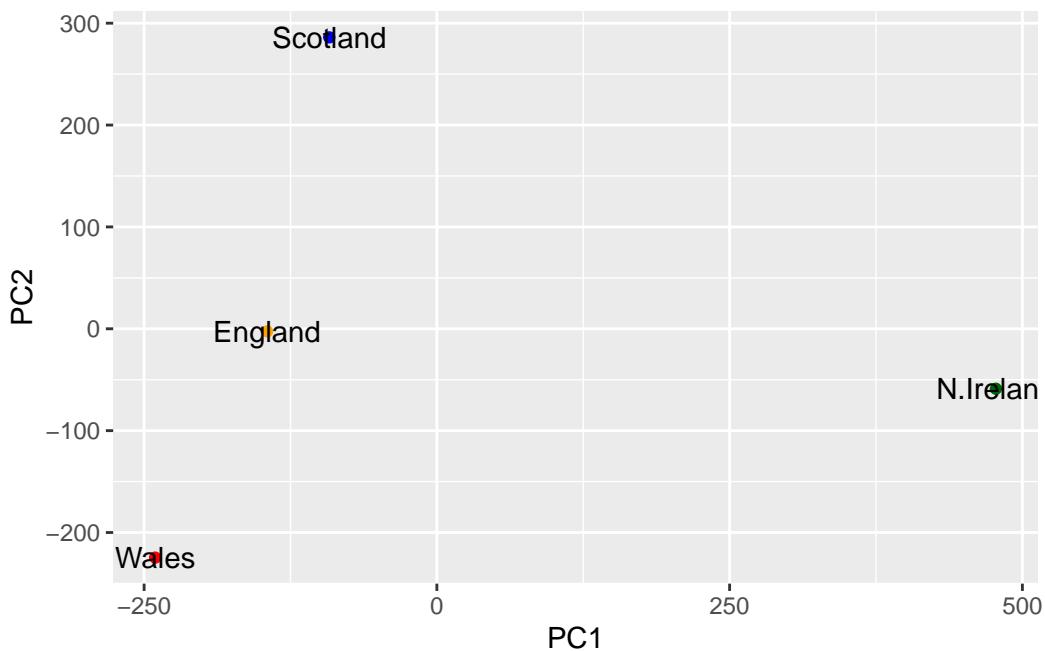
```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
cols <- c("orange", "red", "blue", "darkgreen")
ggplot(pca$x) +
  aes(PC1, PC2, label=rownames(pca$x)) +
  geom_point(col = cols) +
  geom_text()
```



Another major result out of PCA is the so called “variable loadings” or \$rotation that tells us how the original variables (foods) contribute to PCs (the new axis), 0 is the mean the right and left side correspond to ggplot(pca\$x)

```
ggplot(pca$rotation) +  
  aes(PC1, rownames(pca$rotation)) +  
  geom_col()
```

