

Class08: Breast Cancer Analysis Project

Bobbie Morales A15443382

Table of contents

Background	1
Data import	1
Principal Component Analysis	2
Variance explained	6
Communicating PCA results	8
Hierarchical clustering	9
combining PCA and Clustering	10

Background

The goal of this mini-project is for you to explore a complete analysis using the unsupervised learning techniques covered in the last class. We will extend what you've learned by combining PCA as a preprocessing step to clustering using data that consist of measurements of cell nuclei of human breast masses.

The data itself comes from the Wisconsin Breast Cancer Diagnostic Data Set first reported by K. P. Benne and O. L. Mangasarian: "Robust Linear Programming Discrimination of Two Linearly Inseparable Sets".

Values in this data set describe characteristics of the cell nuclei present in digitized images of a fine needle aspiration (FNA) of a breast mass

Data import

```
fna.data <- "WisconsinCancer (2).csv"  
wisc.df <- read.csv(fna.data, row.names=1)
```

Make sure we do not include sample id or diagnosis columns in the data that we analyze below

```
diagnosis <- as.factor(wisc.df$diagnosis)
wisc.data <- wisc.df[,-1]
dim(wisc.data)
```

```
[1] 569 30
```

##Exploratory data analysis >Q1. Q1. How many observations are in this dataset?

There are 569 observations/samples/patients in the data set >Q2. How many of the observations have a malignant diagnosis?

```
sum(wisc.df$diagnosis == "M")
```

```
[1] 212
```

```
table(wisc.df$diagnosis)
```

	B	M
357	212	

Q3. How many variables/features in the data are suffixed with _mean?

```
#colnames(wisc.data)
length(grep("_mean", colnames(wisc.data)))
```

```
[1] 10
```

Principal Component Analysis

The main function in base R for PCA is called `prcomp()`. A optional argument `scale` should nearly always be switched to `scale = TRUE` for this function


```
wisc.pr <- prcomp(wisc.data, scale = TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Let's make our main result figure - the "PC plot" or "score plot" "coordination plot" ..

```
library(ggplot2)
ggplot(wisc.pr$x, aes(x = PC1, y = PC2, col = diagnosis)) +
  geom_point()
```



Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

0.4427 is captured by PC1, as seen by the proportion variance row.

```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005

```

Cumulative Proportion  0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
PC29      PC30
Standard deviation      0.02736 0.01153
Proportion of Variance 0.00002 0.00000
Cumulative Proportion  1.00000 1.00000

```

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?


Three principal components are required, we see PC3 has a cumulative proportion of 0.72636.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

Seven principal components are required, we see PC7 has a cumulative proportion of 0.91010.

Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

```
biplot(wisc.pr)
```




the plot in question 7 is crowded, we cannot see specific data points. We can not draw any conclusions from this plot.

Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

we can see 2 distinct clusters when we plot PC1 vs PC3. We can see that PC1 is separating malignant and benign. Compared to PC1 and PC2, this plot is shifted down, indicating that PC3 might have less of an affect. In PC1 vs PC2, there is a clear elbow point where we could cut the data to show two distinct groups.

```
ggplot(wisc.pr$x, aes(col = diagnosis ,  
x = PC1, y = PC3)) +  
geom_point()
```



Variance explained


Below is the same as the variance explained above. This is how it would be done manually.

```
pr.var <- wisc.pr$sdev^2  
pve <- pr.var / sum(pr.var)  
pve
```


```
[1] 4.427203e-01 1.897118e-01 9.393163e-02 6.602135e-02 5.495768e-02
```

```
[6] 4.024522e-02 2.250734e-02 1.588724e-02 1.389649e-02 1.168978e-02  
[11] 9.797190e-03 8.705379e-03 8.045250e-03 5.233657e-03 3.137832e-03  
[16] 2.662093e-03 1.979968e-03 1.753959e-03 1.649253e-03 1.038647e-03  
[21] 9.990965e-04 9.146468e-04 8.113613e-04 6.018336e-04 5.160424e-04  
[26] 2.725880e-04 2.300155e-04 5.297793e-05 2.496010e-05 4.434827e-06
```

```
plot(pve, xlab = "Principal Component",  
      ylab = "Proportion of Variance Explained",  
      ylim = c(0, 1), type = "o")
```



```
barplot(pve, ylab = "Percent of Variance Explained",  
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)  
axis(2, at=pve, labels=round(pve,2)*100 )
```



Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC.

The component of the loading vector is -0.26085376.

```
wisc.pr$rotation[,1]
```

radius_mean	texture_mean	perimeter_mean
-0.21890244	-0.10372458	-0.22753729
area_mean	smoothness_mean	compactness_mean
-0.22099499	-0.14258969	-0.23928535
concavity_mean	concave.points_mean	symmetry_mean
-0.25840048	-0.26085376	-0.13816696
fractal_dimension_mean	radius_se	texture_se
-0.06436335	-0.20597878	-0.01742803
perimeter_se	area_se	smoothness_se
-0.21132592	-0.20286964	-0.01453145
compactness_se	concavity_se	concave.points_se
-0.17039345	-0.15358979	-0.18341740

symmetry_se	fractal_dimension_se	radius_worst
-0.04249842	-0.10256832	-0.22799663
texture_worst	perimeter_worst	area_worst
-0.10446933	-0.23663968	-0.22487053
smoothness_worst	compactness_worst	concavity_worst
-0.12795256	-0.21009588	-0.22876753
concave.points_worst	symmetry_worst	fractal_dimension_worst
-0.25088597	-0.12290456	-0.13178394

Hierarchical clustering

scaling data


```
data.scaled <- scale(wisc.data)
```

calculating distance btw all pairs of observations in new scaled data set

```
data.dist <- dist(data.scaled)
```

Create a hierarchical clustering model using complete linkage.

```
wisc.hclust <- hclust(data.dist)
plot(wisc.hclust)
```



```
data.dist
hclust (*, "complete")
```

Results of hierarchical clustering ^~


```
wisc.hclust.clusters <- cutree(wisc.hclust, k=4)
```

combining PCA and Clustering

Q12 Which method gives your favorite results for the same data.dist dataset? Explain your reasoning

The ward.D2 method gave the best results for the data since it was more clear and it allowed us to visualize 2 distinct groups and where they would be divided.

```
d <- dist( wisc.pr$x[,1:7] )
wisc.pr.hclust <- hclust(d, method="ward.D2")
plot(wisc.pr.hclust)
abline(h=70, col="red")
```



Get my cluster membership vector

```
grps<- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
  1   2
216 353
```

Q13. How well does the newly created model with four clusters separate out the two diagnoses?

new hierarchical clustering model did a decent job separating benign and malignant samples with 2 clusters. Cluster 1 has 28 benign samples mixed in with malignant while cluster 2 has 24 malignant samples mixed in with benign. This would mean that they both have ~25 false results, which would result in giving a patient the wrong results. with 4 clusters there is a lot of variance in the malignant results.

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=4)
table(wisc.pr.hclust.clusters, diagnosis)
```

```
           diagnosis
wisc.pr.hclust.clusters   B   M
                         1   0  45
                         2   2  77
                         3  26  66
                         4 329  24
```

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
table(wisc.pr.hclust.clusters, diagnosis)
```

```
           diagnosis
wisc.pr.hclust.clusters   B   M
                         1  28 188
                         2 329  24
```

actual diagnoses

```
table(diagnosis)
```

```
diagnosis
  B   M
357 212
```

14 How well do the hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses

when k=4, Cluster 1 contains the majority of the malignant samples while cluster 3 contains the benign samples. The pca + clustering analysis did a better job of separating the diagnoses.

```
table(wisc.hclust.clusters, diagnosis)
```

wisc.hclust.clusters	diagnosis	
	B	M
1	12	165
2	2	5
3	343	40
4	0	2

#sensitivity TP: 179 FP: 24

sensitivity: $TP / (TP + FN)$

section 4 answer Q14, (much better)